

# File-System Interface Implementation

Yohandi [SID: 120040025]

## Background

Data can be stored on a variety of computer-accessible media, including magnetic disks, magnetic tapes, and optical disks. To make the computer system user-friendly, the operating system presents a consistent logical picture of all data. To construct a logical storage unit, the file, the operating system must first abstract from the physical characteristics of its storage devices. The OS physically stores files on whatever the user specifies. Their data remains unchanged even after a system restart, as their nature is nonvolatile.

For all practical purposes, a file only exists in the abstract. Files are best defined by considering the many tasks that may be done on them. The OS makes system calls available to make it possible to make new files, read existing ones, move files around, delete files, and truncate existing ones. Here, we'll look at the OS requirements for executing these six fundamental file operations. Other related procedures, like renaming a file, should become intuitively obvious.

## Program Design

In this task, we are asked to simulate six file management operations: open a file, write to a file, read from a file, list all available files listed by updated time, display all accessible files sorted by file size, and delete a file. In addition to those six activities, we also developed nine utilities to ensure they run well. We need to create an FCB design that will be the basis of our implementation if we want all of the features to perform as intended. All of our file metadata is kept in a database called FCB. Finally, after much deliberation, we settled on the following FCB layout.

For the FCB itself, we are free to use it however we like. Each FCB size is fixed to 32 bytes. Since there are at most 1024 files, we are expected to use a total of  $1024 \times 32$  bytes from the `volume` array in the File System. For a file, we need to store the filename (which consists of 20 bytes), the time when it was created (which consists of 4 bytes =  $256^4$  possible values), the time when it was last modified (which consists of 4 bytes =  $256^4$  possible values), the size of the file (which consists of 2 bytes =  $256^2$  possible values), and the starting position for the file in the storage block (which consists of 2 bytes =  $256^2$  possible values). Adding those numbers, we have  $20 + 4 + 4 + 2 + 2 = 32$ , which is our exact number.

The six file management operations are included in the five main functions of the program. Those are: `fs_open()`, `fs_read()`, `fs_write()`, `fs_gsys()` for displaying the outputs, and `fs_gsys()` for deleting the file.

- `fs_open()` This procedure is in charge of performing the open operation

on the file system and returning a pointer to the memory location of the opened file. The function first verifies the existence of the requested file name in the FCB section of the volume buffer. To save the entry, simply set the existing flag and update the record if the file name discovered matches the desired one. In such circumstances, there are two potential scenarios to think about: the current file already exists, and the current file does not exist. Two steps are necessary to create a file if the current file does not exist. First, space in the file system must be found for the file. Second, an entry for the new file must be made in the directory.

- **fs\_read()** The contents of the requested address are simply copied into the output buffer, making read operations very straightforward. The size parameter's byte count determines how many items are taken. We use a system call to tell the computer which file to read and where in memory to store the data from the file as it is read. Once more, the directory is combed for the corresponding item, and a read pointer is maintained so the system always knows where it stands concerning the file. After a successful read, the read pointer is modified. Due to the binary nature of most processes' file operations, the current operation location may be stored as a current file-position pointer for each process. This single pointer is used for reading and writing, simplifying the system, and saving on storage.
- **fs\_write()** In the first place, we have to make sure that any preexisting files have been removed. Therefore, we change the value of every element in the requested chunk to 0 at the very beginning of the function's source code (empty). Finally, we use the **fp** parameter to determine how the input is stored in memory (pointer from an open operation). The bitmap is then modified to 1. To save data to a file, we issue a system call and provide it the file's name and the data to save. The computer looks in the specified directory for a file with the given name. The system is responsible for maintaining a write pointer that points to the upcoming write position in the file. Every time a write is made, the write pointer must be incremented.
- **fs\_gsys(LS\_D / LS\_S)** When the parameter is set to **LS\_D**, the function is expected to list all file names in the directory and order by modified time of files. On the other hand, when the parameter is set to **LS\_S**, the function is expected to list all file names and sizes in the directory and order by size. We first keep track of each file entry that exists. Once the correct entry is located in the directory, the pointer to the active file is advanced to the specified offset. Moving data inside a file can be done without requiring any I/O operations. A seek is another name for this action on a file. In addition, we can use a standard bubble sort algorithm (as it is easy to implement) to sort the file before displaying it to the user. The detail for the comparison is not necessary to be discussed; however, one vital key to be noted is that we need to involve a combination of the logic operations. The  $O(N^2)$  bubble sort algorithm works pretty well because  $N$ , which

denotes the number of files, is less than or equal to 1024. Of course, an algorithm such as merge sort exists and provides  $O(N\log N)$  computation; however, it is only necessary when  $N$  is relatively large enough, such as  $N \geq 10^5$ .

- **fs\_gsys(RM, filename)** When a file needs to be removed, it is first looked for in the directory structure. After locating the matching directory entry, we delete it and make the freed space available for use by other files. As the implementation, we first scan the FCB and compare it to the current configuration. Whenever a file with the specified name is retrieved, the corresponding index should be saved. After determining the index of the desired entry, we set the bitmap to 1 and replace all occurrences of 0 in the FCB and storage.

## Program Execution

### Program Environment

Note that to execute the program, it is already assumed that the user has already set up access to the school's computation resources.

### Login to Cluster

- Open a new terminal and type `ssh {Student_ID}@CSC4005_cluster` to connect to the remote cluster login with SSH protocol.
- Enter a password to log in.

### Transfer Files to Cluster

- Open a new terminal and locate the files that will be transferred.
- Type `scp {file} {Student_ID}@CSC4005_cluster:.` Here, a file can be in the format of `FILENAME.zip` if multiple files are to be sent. This will send files directly to `nfsmnt/{Student_ID}` destination.

### Execution Steps

- After login in to the cluster, extract the zip files if the files are zipped.
- It is required that in the destination, there are exactly 6 files: `data.bin`, `main.cu`, `slurm.sh`, `user_program.cu`, `file_system.cu`, and `file_system.h`.
- To submit a batch, type `sbatch ./slurm.sh`. After execution is done, new files named `result.out` and `snapshot.bin` will be printed in the same destination. `result.out` shows the compile and run result.
- Another method is by typing both `nvcc --relocatable-device-code=true main.cu user_program.cu virtual_memory.cu -o test` to directly compile the CUDA script using the `nvcc` compiler and `srunk ./test` to run the compiled execution file. Note that this step produces the same

result as the previous step; however, `result.out` will not be printed as the content will be printed directly to the terminal.

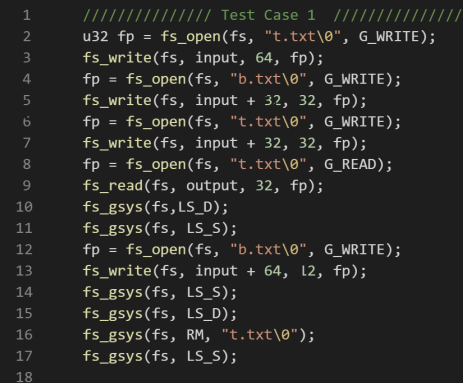
### Transfer Files from Cluster

- If the current terminal is still in the cluster root, type `exit` to exit from the cluster.
- Type `scp 120040025@CSC4005_cluster:~/file} .` to retrieve the file back to the current folder.

### Result

When changing a test case, kindly edit the `user_program.cu` file and transfer it back to the cluster.

#### Test Case #1



```
1  ////////////// Test Case 1 //////////////
2  u32 fp = fs_open(fs, "t.txt\0", G_WRITE);
3  fs_write(fs, input, 64, fp);
4  fp = fs_open(fs, "b.txt\0", G_WRITE);
5  fs_write(fs, input + 32, 32, fp);
6  fp = fs_open(fs, "t.txt\0", G_WRITE);
7  fs_write(fs, input + 32, 32, fp);
8  fp = fs_open(fs, "t.txt\0", G_READ);
9  fs_read(fs, output, 32, fp);
10 fs_gsys(fs, LS_D);
11 fs_gsys(fs, LS_S);
12 fp = fs_open(fs, "b.txt\0", G_WRITE);
13 fs_write(fs, input + 64, 12, fp);
14 fs_gsys(fs, LS_S);
15 fs_gsys(fs, LS_D);
16 fs_gsys(fs, RM, "t.txt\0");
17 fs_gsys(fs, LS_S);
18
```

In this case, we first create a file named `t.txt` and write content with a size of 64 bytes from the first 64 bytes in `data.bin`. Then, we create another file named `b.txt` and write content with a size of 32 bytes from `[32..63]` in `data.bin`. We do the exact same previous operation for `t.txt`.

After the first 3 writing cases, the program will then read the value from the first 32 bytes in `t.txt` and write it to `snapshot.bin`. This is the reason why `snapshot.bin` is filled with 32 bytes from `[32..63]` in `data.bin`.

snapshot\_1.bin

HEX

offset

0

absolute

relative

address	00	01	02	03	04	05	06	07	08	09	10	11	12	13	14	15	Ascll	<input type="checkbox"/> unsigned	<input type="checkbox"/> bigendian
00000000	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	0000000000000000		
00000010	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	6f	0000000000000000		
00000020	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		
00000030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		
00000040	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		
00000050	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		
00000060	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		
00000070	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		
00000080	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		
00000090	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		
000000a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		
000000b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		
000000c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		
000000d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		
000000e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		
000000f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....		

snapshot.bin

After the reading case, the program will then list the existing files sorted by two states: `LS_D` and `LS_S`. With a `LS_D` state, it will list the files sorted by their last modified time. This will print `t.txt` then `b.txt` since `t.txt` was the last modified version. With a `LS_S` state, it will list the files sorted by their file size. This will print `t.txt 32` and then `b.txt 32` since `t.txt` was the first one to be created.

The program will then write the content of 12 bytes from `[64..75]` in `data.bin`.

After the writing case, the program will again list the existing files sorted by two states: `LS_D` and `LS_S`. With a `LS_S` state, it will list the files sorted by their file size. This will print `t.txt 32` then `b.txt 12` since `t.txt` has a bigger size than `b.txt`. With a `LS_D` state, it will list the files sorted by their last modified time. This will print `b.txt` then `t.txt` since `b.txt` was the last modified version.

The program will remove the `t.txt` file and list the existing files with `LS_S` state. This will print `b.txt 12` since `b.txt` is the only file.

```
result_1.out x
result_1.out
77  ===sort by modified time===
78  t.txt
79  b.txt
80  ===sort by file size===
81  t.txt 32
82  b.txt 32
83  ===sort by file size===
84  t.txt 32
85  b.txt 12
86  ===sort by modified time===
87  b.txt
88  t.txt
89  ===sort by file size===
90  b.txt 12
91
```

result.out

## Test Case #2

```
1  ////////////// Test Case 2 //////////////
2  u32 fp = fs_open(fs, "t.txt\0", G_WRITE);
3  fs_write(fs,input, 64, fp);
4  fp = fs_open(fs,"b.txt\0", G_WRITE);
5  fs_write(fs,input + 32, 32, fp);
6  fp = fs_open(fs,"t.txt\0", G_WRITE);
7  fs_write(fs,input + 32, 32, fp);
8  fp = fs_open(fs,"t.txt\0", G_READ);
9  fs_read(fs,output, 32, fp);
10 fs_gsys(fs,LS_D);
11 fs_gsys(fs,LS_S);
12 fp = fs_open(fs,"b.txt\0", G_WRITE);
13 fs_write(fs,input + 64, 12, fp);
14 fs_gsys(fs,LS_S);
15 fs_gsys(fs,LS_D);
16 fs_gsys(fs,RM, "t.txt\0");
17 ...
18
```

snapshot\_2.bin

offset 0absoluterelative

address

00 01 02 03 04 05 06 07 08 09 10 11 12 13 14 15

Ascll

■ unsigned

■ bigendian

00000000	6f 6f 6f 6f 6f 6f 6f 6f	6f 6f 6f 6f 6f 6f 6f 6f	0000000000000000
00000010	6f 6f 6f 6f 6f 6f 6f 6f	6f 6f 6f 6f 6f 6f 6f 6f	0000000000000000
00000020	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....
00000030	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....
00000040	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....
00000050	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....
00000060	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....
00000070	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....
00000080	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....
00000090	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....
000000a0	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....
000000b0	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....
000000c0	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....
000000d0	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....
000000e0	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....
000000f0	00 00 00 00 00 00 00 00	00 00 00 00 00 00 00 00	.....

snapshot.bin

result\_2.out

result\_2.out

```
77 ===sort by modified time===
78 t.txt
79 b.txt
80 ===sort by file size===
81 t.txt 32
82 b.txt 32
83 ===sort by file size===
84 t.txt 32
85 b.txt 12
86 ===sort by modified time===
87 b.txt
88 t.txt
89 ===sort by file size===
90 b.txt 12
91 ===sort by file size===
92 *ABCDEFGHJKLMNPQR 33
93 )ABCDEFGHJKLMNPQR 32
94 (ABCDEFGHJKLMNPQR 31
95 *ABCDEFGHJKLMNPQR 30
```

result.out





```
result_3.out x
result_3.out
81  ===sort by modified time===
82  t.txt
83  b.txt
84  ===sort by file size===
85  t.txt 32
86  b.txt 32
87  ===sort by file size===
88  t.txt 32
89  b.txt 12
90  ===sort by modified time===
91  b.txt
92  t.txt
93  ===sort by file size===
94  b.txt 12
95  ===sort by file size===
96  *ABCDEFGHIIJKLMNOPQR 33
97  )ABCDEFGHIIJKLMNOPQR 32
98  (ABCDEFGHIIJKLMNOPQR 31
99  'ABCDEFGHIIJKLMNOPQR 30
```

result.out

#### Test Case #4

```
1  ////////////// Test Case 4 //////////////
2  u32 fp = fs_open(fs, "32-block-0", G_WRITE);
3  fs_write(fs, input, 32, fp);
4  for (int j = 0; j < 1023; ++j) {
5      char tag[] = "1024-block-????";
6      int i = j;
7      tag[11] = static_cast<char>(i / 1000 + '0');
8      i = i % 1000;
9      tag[12] = static_cast<char>(i / 100 + '0');
10     i = i % 100;
11     tag[13] = static_cast<char>(i / 10 + '0');
12     i = i % 10;
13     tag[14] = static_cast<char>(i + '0');
14     fp = fs_open(fs, tag, G_WRITE);
15     fs_write(fs, input + j * 1024, 1024, fp);
16 }
17 ...
18
```

```

snapshot_4.bin X
offset 0 absolute relative
address 00 01 02 03 04 05 06 07 08 09 10 11 12 13 14 15 Ascii unsigned bigendian
00000000 64 64 64 64 64 64 64 64 64 64 64 64 64 64 64 64 ddddddddddddddd
00000010 64 64 64 64 64 64 64 64 64 64 64 64 64 64 64 64 ddddddddddddddd
00000020 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f ooooooooooooooooo
00000030 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f 6f ooooooooooooooooo
00000040 63 63 63 63 63 63 63 63 63 63 63 63 aa 13 68 df ccccccccccc.h.
00000050 13 09 57 56 d3 19 c2 70 5b 39 ab 09 1a c2 6f ad ..WV.p[9.o.
00000060 68 52 80 14 ba b6 12 b9 f2 fc da 9d 10 c2 fc 23 hR.....#
00000070 cc d4 f8 20 6d 3b 10 c8 74 3b d1 8f fd c0 bc e5 ...m;.t;....
00000080 13 3d fa 4e f3 0d 08 66 89 e2 83 99 25 81 3c 71 .=N...f...%.<q
00000090 56 35 91 c3 ef 21 0c 65 5c 5d 74 5a 1f 31 41 32 V5...!.e\]tZ1A2
000000a0 6f bb 80 e2 48 08 4a d1 6a cd ea 90 ce 27 02 25 o...H.J.j....'%.%
000000b0 db 14 68 cb 35 74 b0 92 52 25 ec 71 57 ad 23 46 .@h.5t..R%.qW.#F
000000c0 69 a4 29 31 ac f2 03 18 41 ed 28 10 94 2a b5 70 i.)1...@A.(...*.p
000000d0 3e 1e bc f3 13 6d 86 65 93 f2 d6 6a a1 f9 b0 8a >@...m.e...j....
000000e0 1e 59 bc 4b 4d 3f e2 8e 2e 0b 9e c2 b4 d3 b3 73 @Y.KM?.....s
000000f0 f2 70 67 06 5d 6d 6b f0 60 c1 da 81 3b 8b 8c 5a .pg.]mk. ...;..Z

```

First 16 lines of snapshot.bin

```

result_4.out X
result_4.out
65 triggering gc
66 ===sort by modified time===
67 1024-block-1023
68 1024-block-1022
69 1024-block-1021
70 1024-block-1020
71 1024-block-1019
72 1024-block-1018
73 1024-block-1017
74 1024-block-1016
75 1024-block-1015
76 1024-block-1014
77 1024-block-1013
78 1024-block-1012
79 1024-block-1011
80 1024-block-1010
81 1024-block-1009
82 1024-block-1008
83 1024-block-1007

```

result.out

## Reflection

The basics of CUDA programming were our first lesson in this task. Conversely, we picked up the bitmap technique for documenting the files already in memory and the free space between them. Finally, we have a general idea of how to implement the sorting mechanism in the file system when showing the existing files and what happens when the remove operation is asked to delete a file. Considering the significance of the file system handled by the operating system, this project invites students to play the role of the operating system in carrying out certain activities. Finally, it has been demonstrated that the program runs

all test cases successfully.